


# Monitors & Condition Synchronization



# Monitors & Condition Synchronization

---

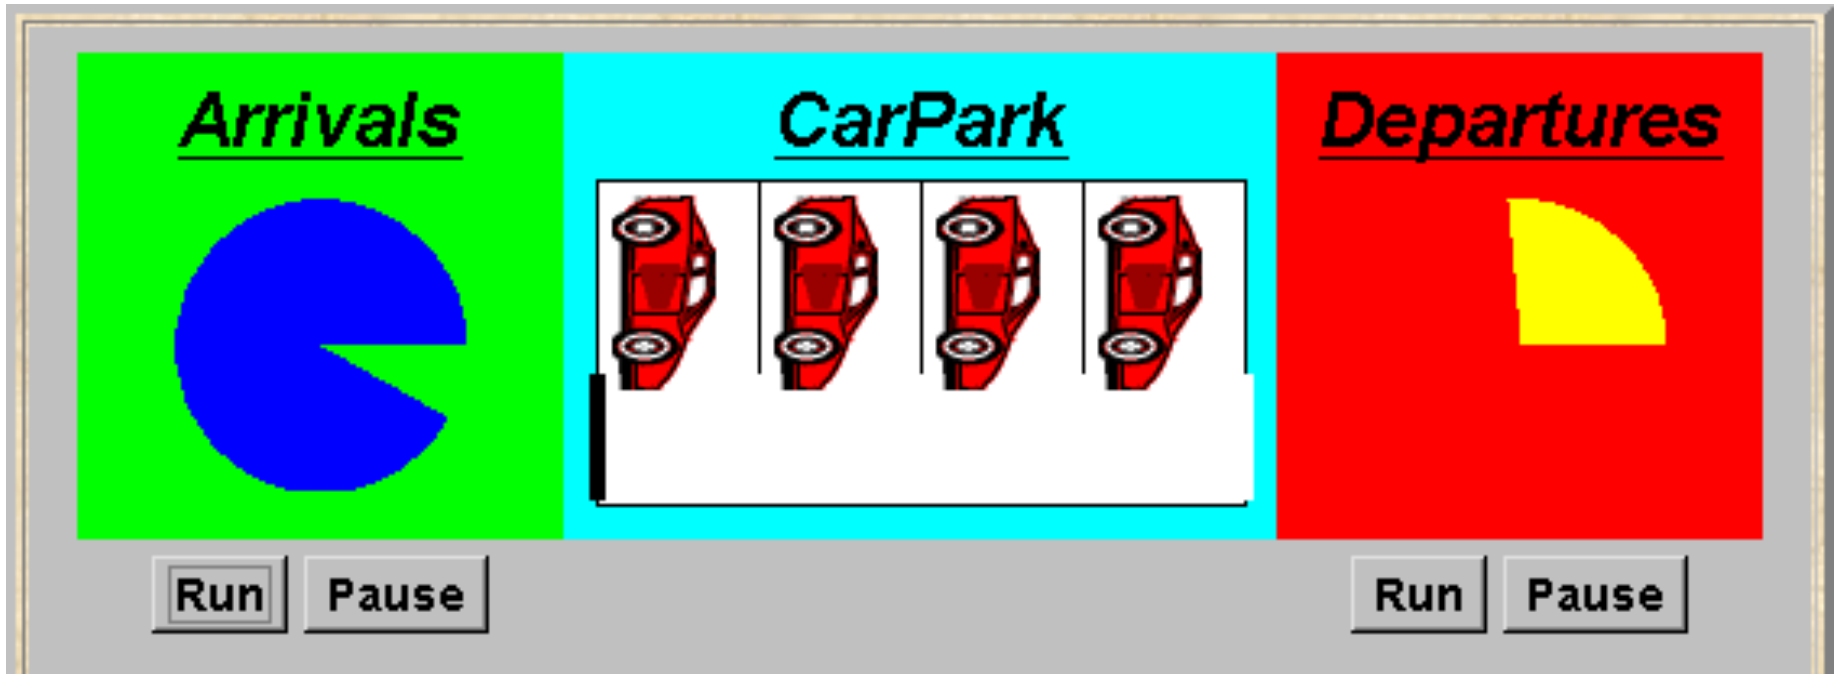
**Concepts:** monitors:

- encapsulated data + access procedures
- mutual exclusion + condition synchronization
- single access procedure active in the monitor
- nested monitors

**Models:** guarded actions

**Practice:** private data and synchronized methods (exclusion).  
`wait()`, `notify()` and `notifyAll()` for condition synch.  
single thread active in the monitor at a time

## 5.1 Condition Synchronization



- A **controller** is required for a carpark, which
  - only permits cars to enter when the carpark is not full
  - not permit a car to leave if there is no car in the carpark.
- Car arrival and departure are simulated by separate threads.

# Carpark Model

---

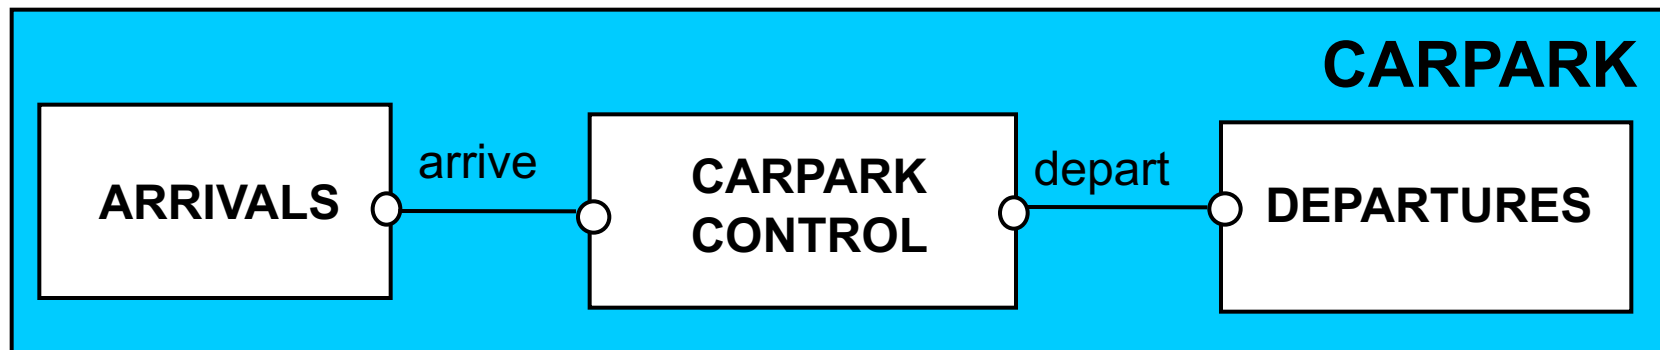
- ◆ Events or actions of interest?

arrive and depart

- ◆ Identify processes.

arrivals, departures and carpark control

- ◆ Define each process and interactions (structure).




## Carpark Model (FSP)

```
CARPARKCONTROL (N=4) = SPACES [N] ,  
SPACES [i:0..N] = (when (i>0) arrive->SPACES [i-1]  
                  | when (i<N) depart->SPACES [i+1]  
                  ) .  
  
ARRIVALS      = (arrive->ARRIVALS) .  
DEPARTURES    = (depart->DEPARTURES) .  
  
|| CARPARK =  
    (ARRIVALS || CARPARKCONTROL (4) || DEPARTURES) .
```

Guarded actions are used  
to control **arrive** and  
**depart**.

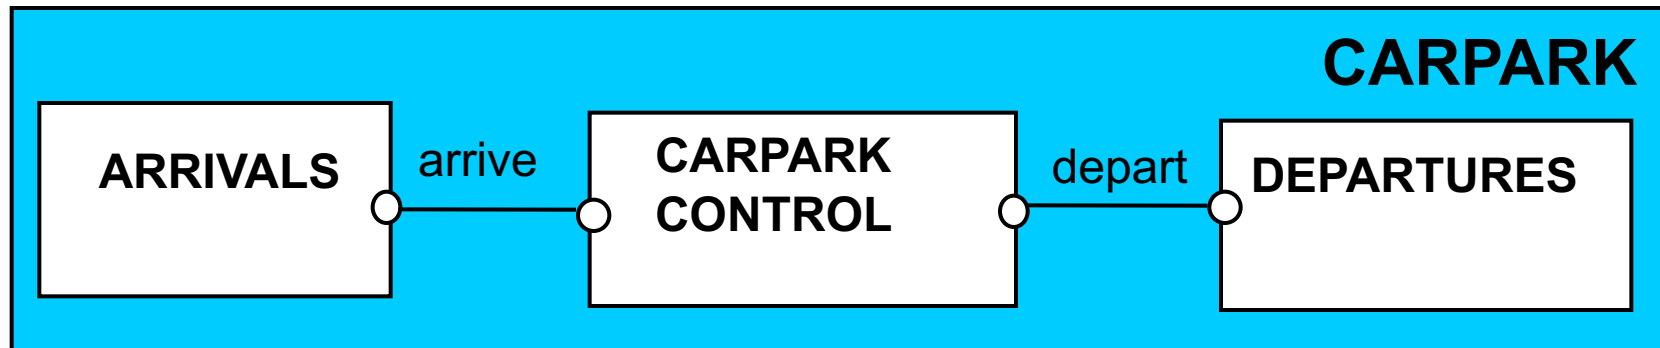
Concurrency: monitors & condition synchronizati



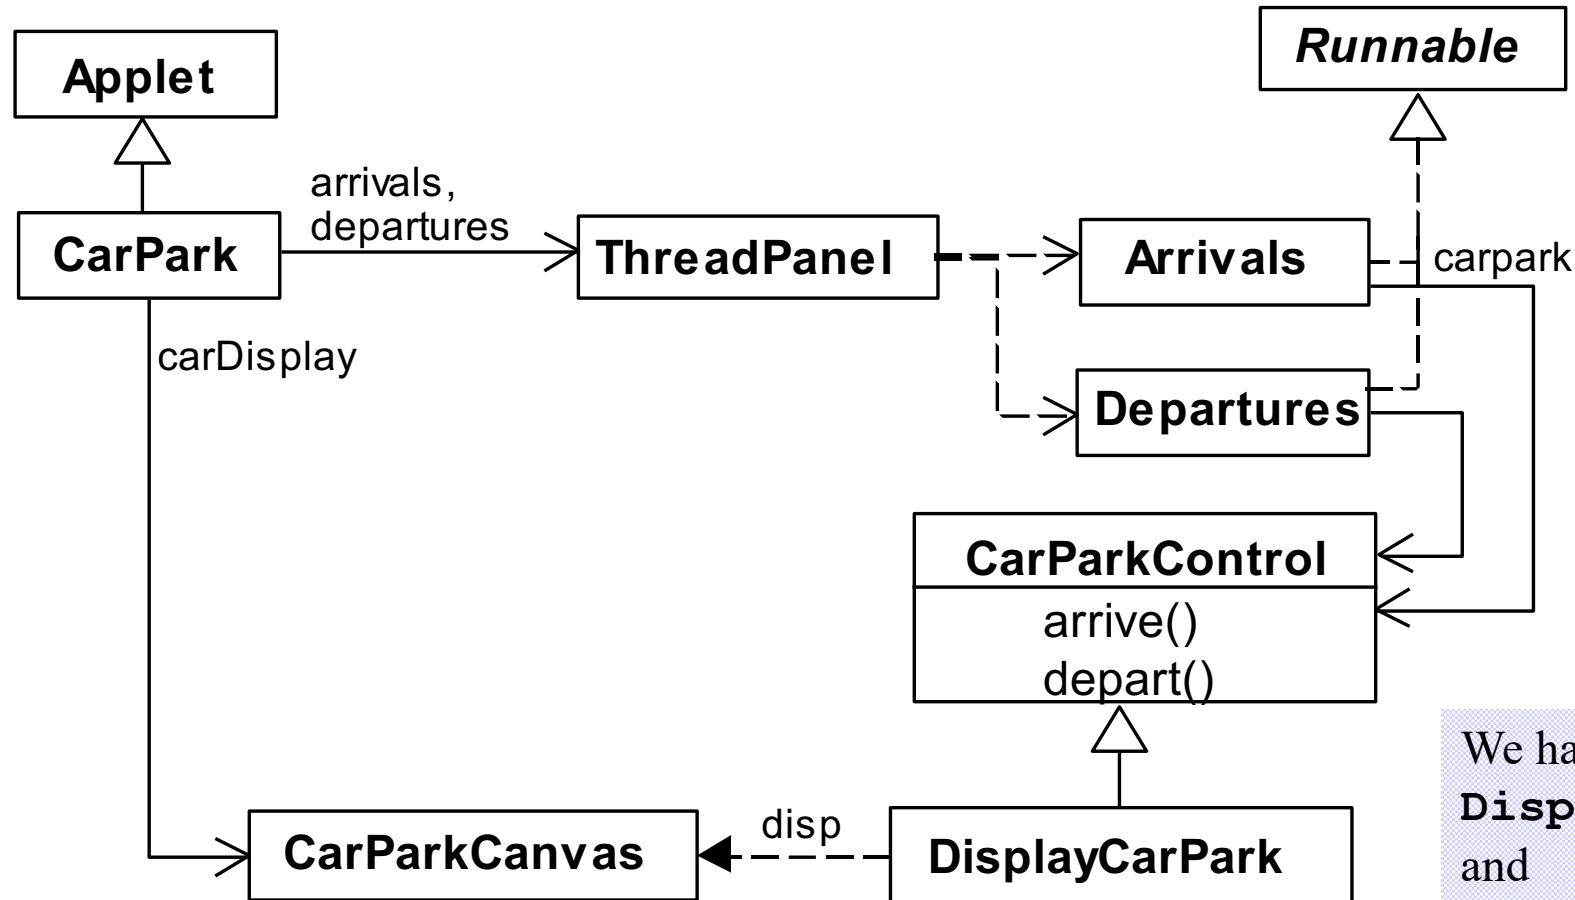
# Carpark Program

---

- ♦ Model - all entities are processes interacting by actions
- ♦ Program - need to identify threads and monitors
  - ♦ thread - active entity which initiates actions—arrive, depart
  - ♦ monitor - passive entity which responds to actions.



# Carpark Program – Class Diagram



We have omitted **DisplayThread** and **GraphicCanvas** threads managed by **ThreadPanel**.

# Carpark Program

---

- **Arrivals** and **Departures** implement **Runnable**
- **CarParkControl** provides the **control-condition synchronization**.
- **Instances** of these are created by the **start()** method of the **CarPark** applet :

```
public void start() {  
    CarParkControl c =  
        new DisplayCarPark(carDisplay, Places) ;  
    arrivals.start(new Arrivals(c)) ;  
    departures.start(new Departures(c)) ;  
}
```



## Carpark Program - Arrivals and Departures threads

---

```
class Arrivals implements Runnable {
    CarParkControl carpark;

    Arrivals(CarParkControl c) {carpark = c;}

    public void run() {
        try {
            while(true) {
                ThreadPanel.rotate(330);
                carpark.arrive();
                ThreadPanel.rotate(30);
            }
        } catch (InterruptedException e) {}
    }
}
```

Similarly Departures  
which calls  
carpark.depart().

# Carpark Program – CarParkControl Monitor

---

```
class CarParkControl {
    protected int spaces;
    protected int capacity;

    CarParkControl(int n)
        {capacity = spaces = n;}

    synchronized void arrive() {
        ... --spaces; ...
    }

    synchronized void depart() {
        ... ++spaces; ...
    } }
}
```

*mutual exclusion by  
synch methods*

*(condition)  
synchronization?*

*block if full?  
(spaces==0)*

*block if empty?  
(spaces==N)*

## Monitor:

- a synchronization construct allowing threads to have both
  - mutual exclusion
  - ability to wait for a certain condition to become true
- consists of a mutex (lock) object and condition variables.

## Condition Synchronization in Java

---

Java provides a **thread wait set** per **monitor** (actually per **object**) with the following methods:

```
public final void wait()
```

```
throws InterruptedException
```

Waits to be notified by another thread. The **waiting thread releases** the **synchronization lock** associated with the **monitor**. When notified, the thread must wait to **reacquire** the monitor before resuming execution.

```
public final void notify()
```

Wakes up a **single thread** that is waiting on this object's wait set.

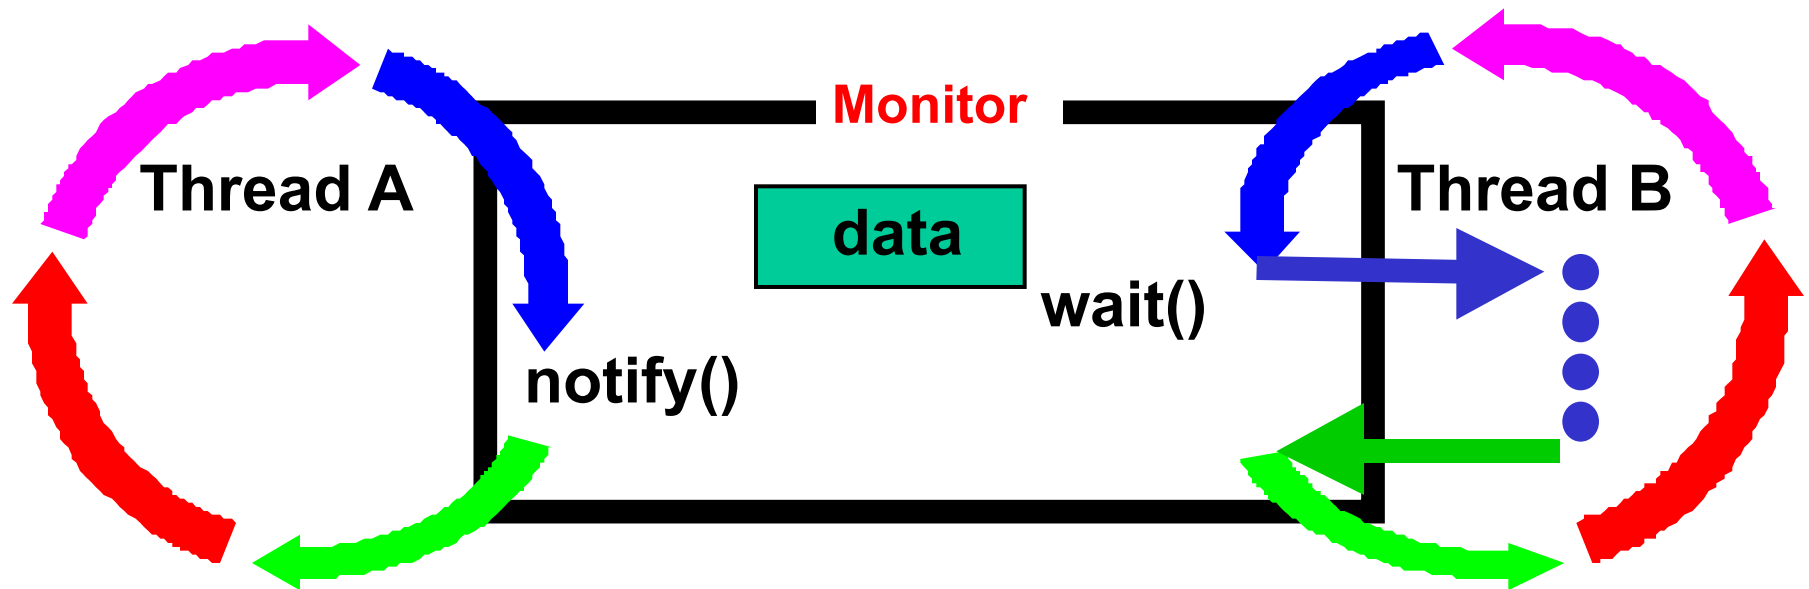
```
public final void notifyAll()
```

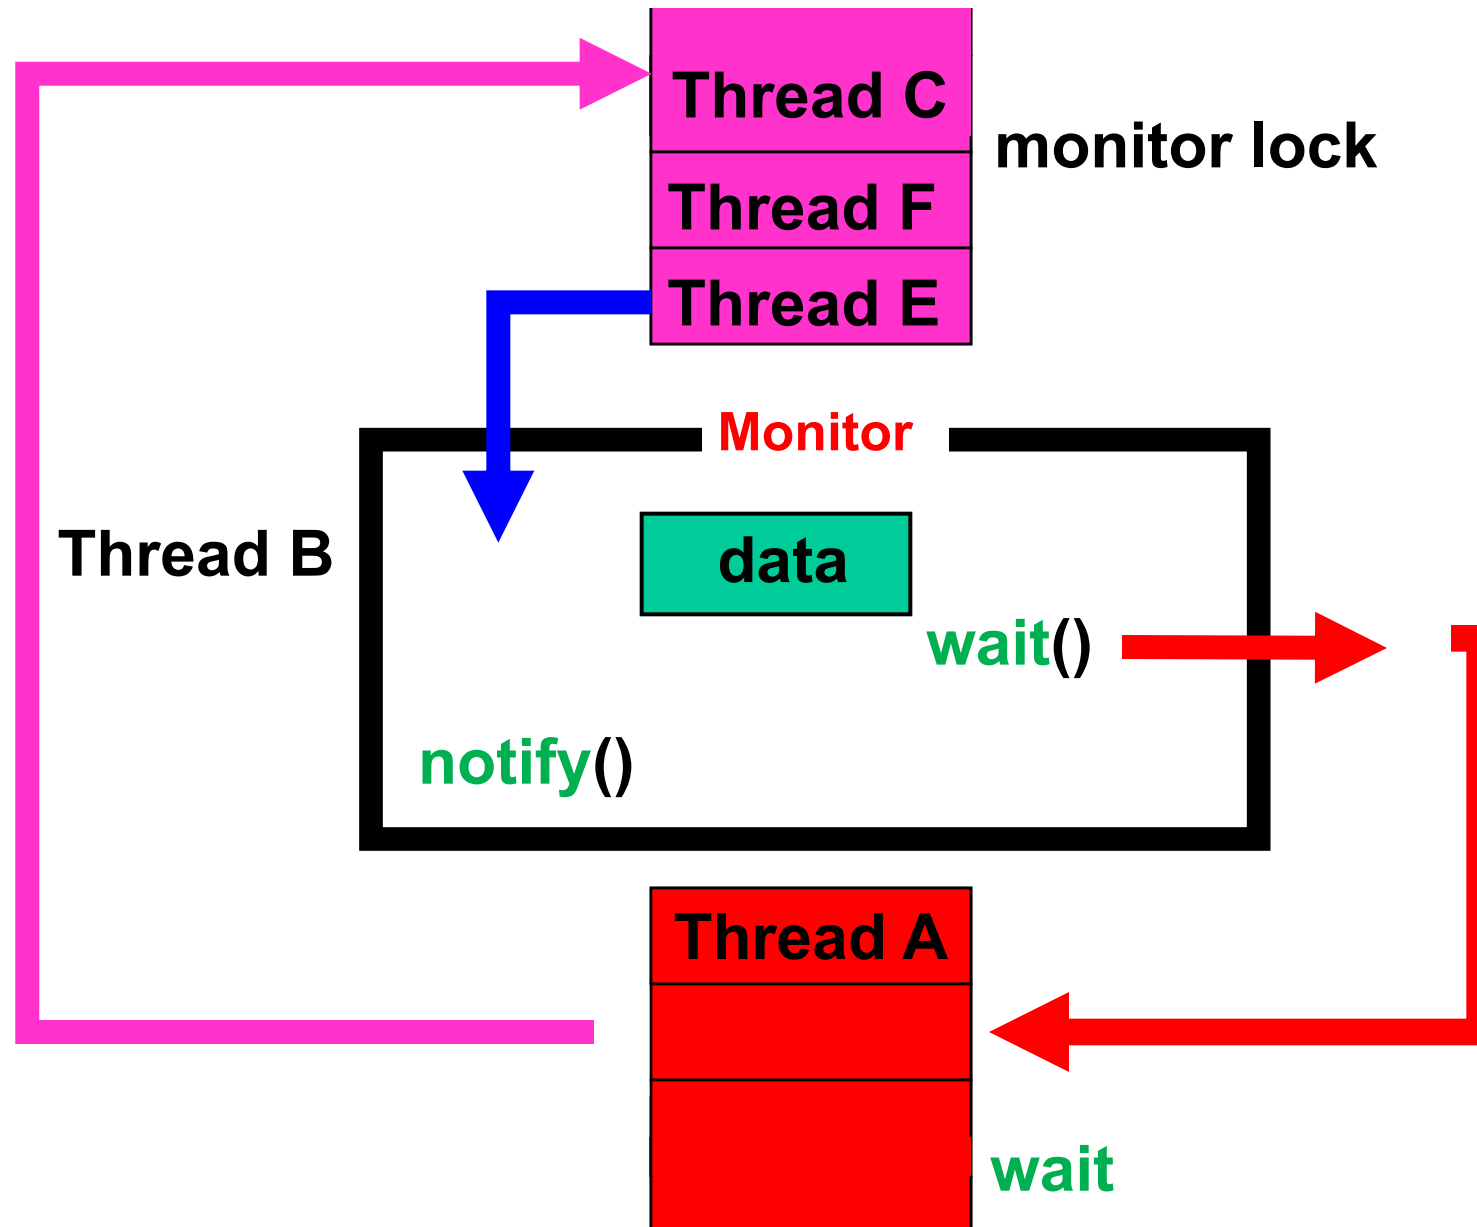
Wakes up **all threads** that are waiting on this object's wait set.

# Condition Synchronization in Java

---

- **Entering a monitor:** when a thread acquires the mutual exclusion **lock** associated with the monitor
- **Exiting the monitor:** when it releases the **lock**.
- **Wait()** - causes the thread to exit the monitor, permitting other threads to enter the monitor.





## Condition Synchronization in Java

---

FSP:     when *cond* act -> NEWSTAT

```
Java:  public synchronized void act()
        throws InterruptedException
    {
        while (!cond) wait();
        // modify monitor data
        notifyAll();
    }
```

The **while** loop is necessary to **retest** the condition *cond* to ensure that *cond* is indeed **satisfied** when it **re-enters** the **monitor**.

**notifyall()** is necessary to **awaken other thread(s)** that may be waiting to enter the monitor now that the monitor data has been **changed**.

## CarParkControl – Condition Synchronization

---

```
class CarParkControl {
    protected int spaces;
    protected int capacity;

    CarParkControl(int n)
        {capacity = spaces = n;}

    synchronized void arrive() throws InterruptedException {
        while (spaces==0) wait();
        --spaces;
        notifyAll();
    }

    synchronized void depart() throws InterruptedException {
        while (spaces==capacity) wait();
        ++spaces;
        notifyAll();
    }
}
```

Note: **notify()** wakes up the **first thread** that called **wait()** on the **same object**.  
**notifyAll()** wakes up **all the threads** that called **wait()** on the **same object**.

## Models of Monitors - Summary

---

**Active entities** (that **initiate actions**) are implemented as **threads**.

**Passive entities** (that **respond to actions**) are implemented as **monitors**.

Each **guarded** action in the model of a monitor is implemented as a **synchronized** method which uses a **while loop** and **wait()** to **implement the guard**.

Changes in the **state** of the monitor are signalled to **waiting threads** using **notify()** or **notifyAll()**.



## 5.2 Semaphores

---

- **Semaphores**: used for **controlling access** to a **common resource** by **multiple processes** in a **concurrent system**.
- **Semaphore  $s$** : **non-negative, integer variable**.
- The only **operations** permitted on  **$s$**  are  **$up(s)$**  and  **$down(s)$** .

**$down(s)$** : if  **$s > 0$**  then  
decrement  **$s$**   
else  
**block execution** of the calling process
- **Blocked processes** are held in a **FIFO queue**.

**$up(s)$** : if processes are blocked on  **$s$**  then  
awaken one of them  
increment  **$s$**

## Modelling Semaphores

---


- Semaphores take a **finite range of values**. If the range is exceeded then **ERROR**.
- **N** with **initial value = 0**.

```
const Max = 3  
range Int = 0..Max
```

```
SEMAPHORE (N=0) = SEMA [N] ,  
SEMA [v: Int]    = (up -> SEMA [v+1]  
                  | when (v>0) down -> SEMA [v-1]  
                  ) ,  
SEMA [Max+1]     = ERROR.
```

# Modelling Semaphores

---



- Action **down** is only **accepted** when value **v** of the semaphore is greater than 0.
- Action **up** is **not guarded**.
- Trace to a violation:  
 $\text{up} \rightarrow \text{up} \rightarrow \text{up} \rightarrow \text{up}$

## Semaphore Demo - Model

---

Three processes  $p[1..3]$  use a shared semaphore *mutex* to ensure *mutually exclusive access* (action *critical*) to some *resource*.

LOOP = (mutex.down -> critical -> mutex.up -> LOOP) .

|| SEMADEMO = ( p[1..3]:LOOP  
|| {p[1..3]}::mutex:SEMAPHORE(1) ) .

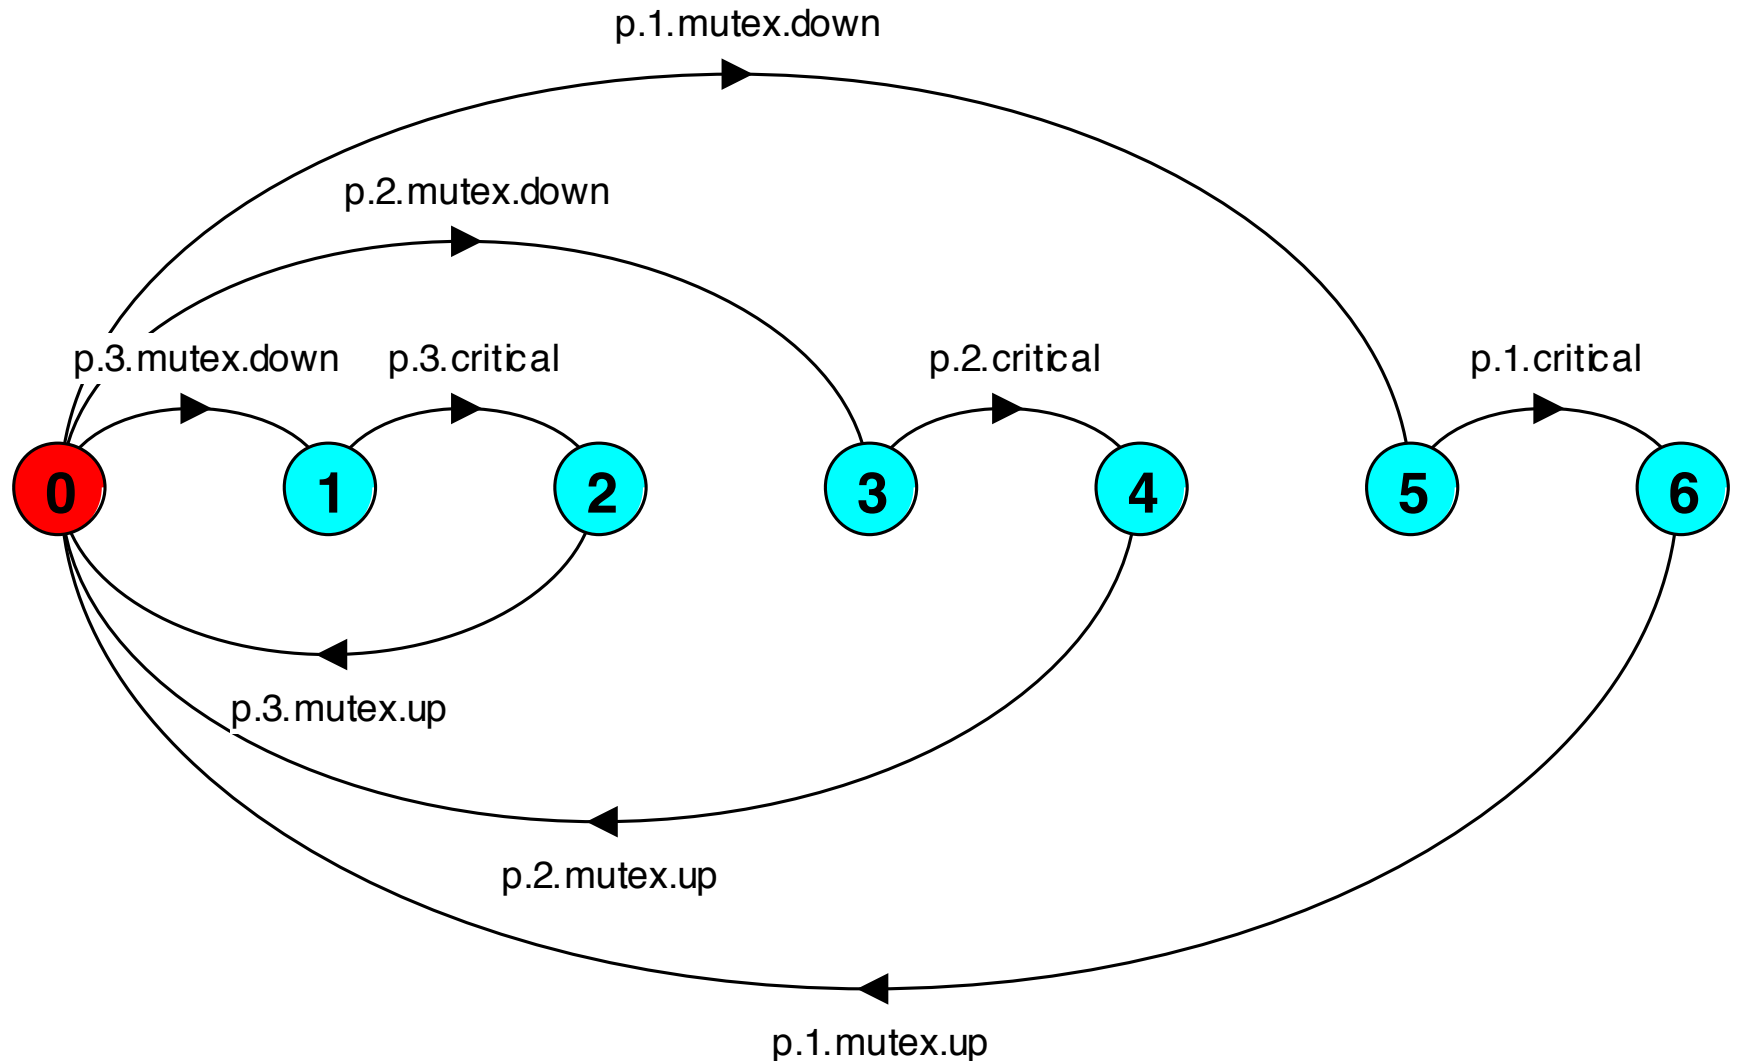
For mutual exclusion, the semaphore initial value is 1. *Why?*

*Is the ERROR state reachable for SEMADEMO?*

*Is a binary semaphore sufficient (i.e. Max=1) ?*

Note: semaphore can be counted, while mutex can only count to 1.

# Semaphore Demo - Model



# Semaphores in Java

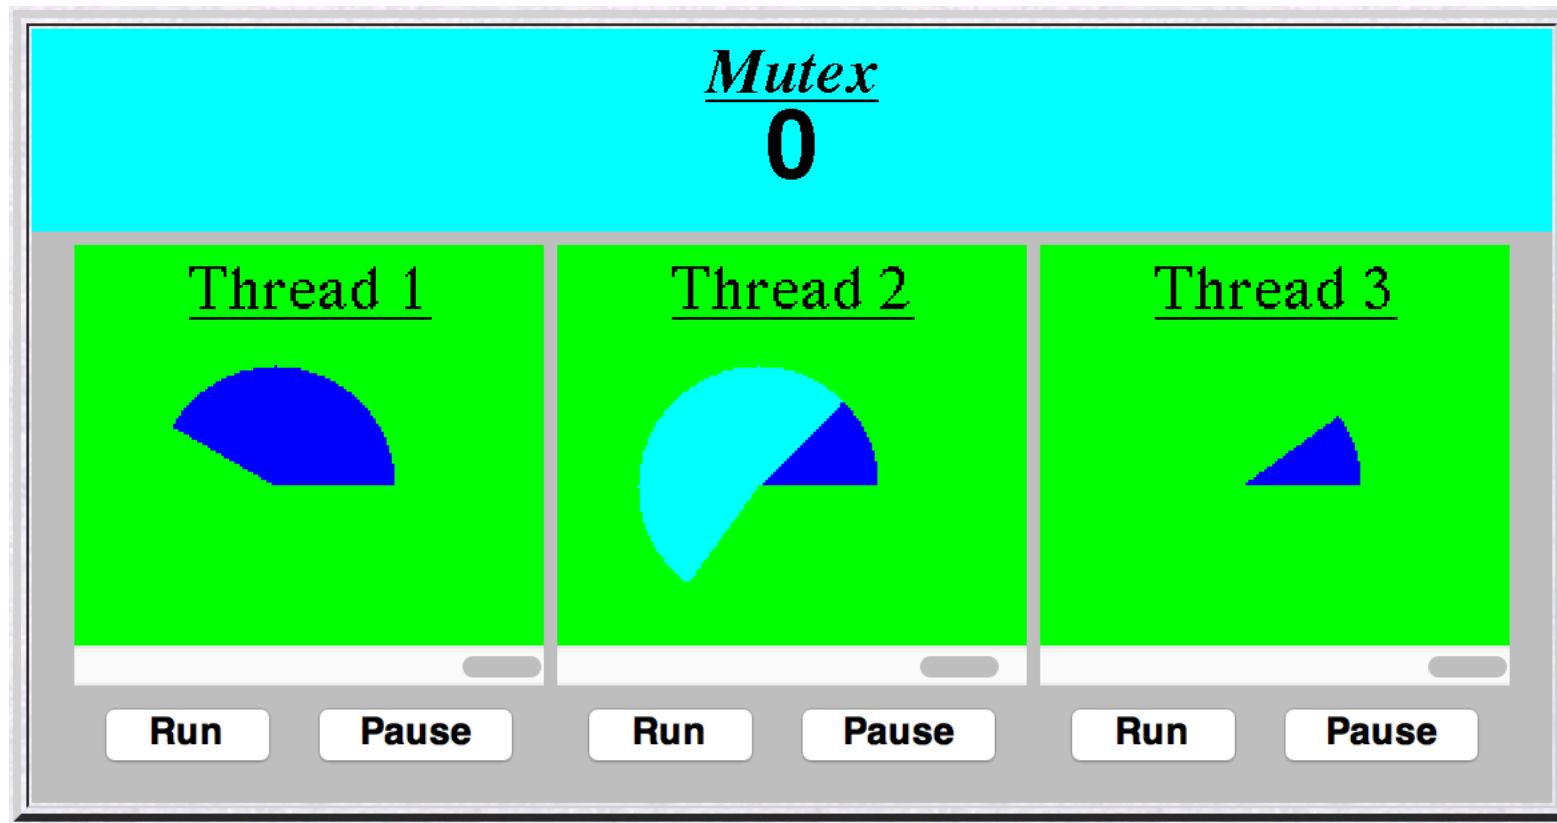
- Semaphores are passive objects, implemented as monitors.

```
public class Semaphore {  
    private int value;  
  
    public Semaphore (int initial)  
        {value = initial;}  
  
    synchronized public void up() {  
        ++value;  
        notifyAll();  
    }  
  
    synchronized public void down()  
        throws InterruptedException {  
        while (value== 0) wait();  
        --value;  
    }  
}
```

*Is it safe to use notify() here rather than notifyAll()?*

## SEMADEMO display

---



**current  
semaphore  
value**

**Thread 2 is  
executing  
critical  
actions.**

**Threads 1, 3  
are blocked  
waiting.**

What if we adjust the *time that each thread spends in its critical section* ?

◆ large resource requirement - *more conflict*?

(e.g., more than 67% of a rotation)?

◆ small resource requirement - *no conflict*?

(e.g., less than 33% of a rotation)?

the time a thread spends in its critical section  
should be kept as short as possible.



## SEMADEMO program - MutexLoop

```
class MutexLoop implements Runnable {
    Semaphore mutex;

    MutexLoop (Semaphore sema) {mutex=sema;}


    public void run() {
        try {
            while(true) {
                while(!ThreadPanel.rotate());
                mutex.down();           // get mutual exclusion
                while(ThreadPanel.rotate()); // critical actions
                mutex.up();             // release mutual exclusion
            }
        } catch (InterruptedException e) {}
    }
}
```

Threads and semaphore are created by the applet `start()` method.

`ThreadPanel.rotate()` returns `false` while executing non-critical actions (blue color) and `true` otherwise.

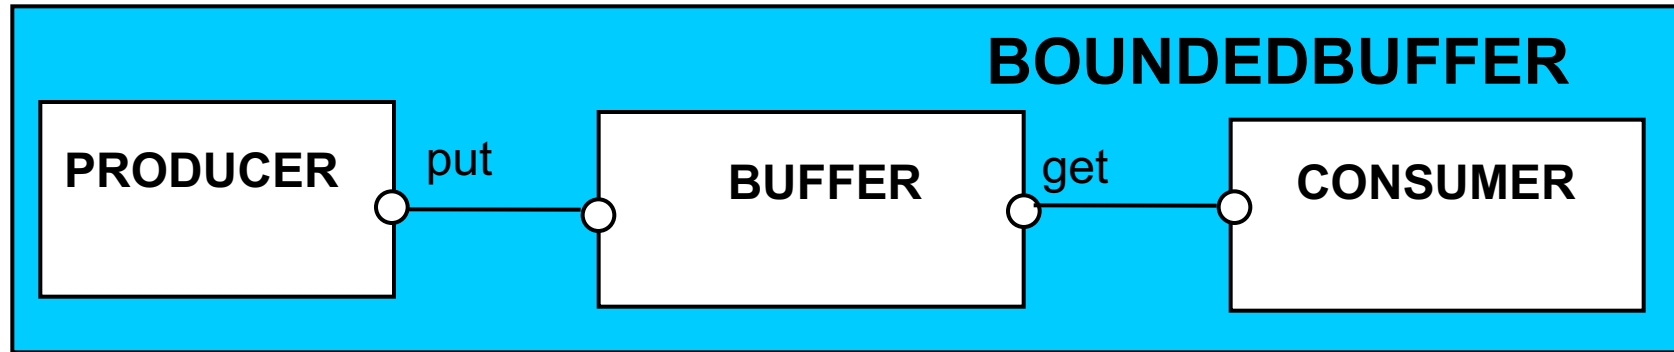
## 5.3 Bounded Buffer

---




- A **bounded buffer** consists of a **fixed number of slots**.
- **Items** are **put** into the buffer by a **producer process** and **removed** by a **consumer process**.

## Bounded Buffer - a Data-Independent Model



The behaviour of BOUNDEDBUFFER is independent of the actual data values, and so can be modelled in a data-independent manner.

LTS:



## Bounded Buffer - a Data-Independent Model

---

```
BUFFER (N=5) = COUNT [ 0 ] ,  
COUNT [ i : 0 .. N ]  
    = (when (i < N) put->COUNT [ i+1 ]  
       | when (i > 0) get->COUNT [ i-1 ]  
       ) .
```

```
PRODUCER = (put->PRODUCER) .  
CONSUMER = (get->CONSUMER) .
```

```
|| BOUNDEDBUFFER =  
(PRODUCER || BUFFER (5) || CONSUMER) .
```

# Bounded Buffer Program – Buffer Monitor

---

```
public interface Buffer <E> {...}

class BufferImpl <E> implements Buffer <E> {
    ...
    public synchronized void put(E o)
        throws InterruptedException {
        while (count == maxSize) wait();
        buf[in] = o; ++count; in=(in+1)%maxSize;
        notifyAll();
    }
    public synchronized E get()
        throws InterruptedException {
        while (count == 0) wait();
        E o = buf[out];
        buf[out]=null; --count; out=(out+1)%maxSize;
        notifyAll();
        return (o);
    }
}
```

# Bounded Buffer Program – Producer Process

---

```
class Producer implements Runnable {
    Buffer buf;
    String alphabet= "abcdefghijklmnopqrstuvwxyz";
    Producer(Buffer b) {buf = b;}

    public void run() {
        try {
            int ai = 0;
            while(true) {
                ThreadPanel.rotate(12);
                buf.put(alphabet.charAt(ai));
                ai=(ai+1) % alphabet.length();
                ThreadPanel.rotate(348);
            }
        } catch (InterruptedException e) {}
    }
}
```

Producer calls  
buf.put().

Consumer calls  
buf.get().

## 5.4 Nested Monitors

---

Instead of using *count* variable and *condition synchronization* directly, use **two semaphores** *full* and *empty* to reflect the **state** of the buffer.

```
class SemaBuffer <E> implements Buffer <E> {  
    ...  
  
    Semaphore full;    // the number of items in Buffer  
    Semaphore empty;   // the number of available spaces in Buffer  
  
    SemaBuffer(int size) {  
        this.size = size; buf =(E[])new Object[size];  
        full = new Semaphore(0);  
        empty= new Semaphore(maxSize); // buffer maxSize  
    }  
    ...  
}
```

## Bounded Buffer Program – Buffer Monitor

```
public interface Buffer <E> {...}

class BufferImpl <E> implements Buffer <E> {
    ...
    public synchronized void put(E o)
        throws InterruptedException {
        while (count == maxSize) wait();
        buf[in] = o; ++count; in=(in+1)%maxSize;
        notifyAll();
    }
    public synchronized E get()
        throws InterruptedException {
        while (count == 0) wait();
        E o = buf[out];
        buf[out]=null; --count; out=(out+1)%maxSize;
        notifyAll();
        return (o);
    }
}
```

*condition variable:*  
*count*



## Nested Monitors - Bounded Buffer Program

```
synchronized public void put(E o)
    throws InterruptedException {
    empty.down();
    buf[in] = o;
    ++count; in=(in+1)% maxSize;
    full.up();
}

synchronized public E get()
    throws InterruptedException{
    full.down();
    E o =buf[out]; buf[out]=null;
    --count; out=(out+1)% maxSize;
    empty.up();
    return (o);
}
```

*Does this behave  
as desired?*

- *empty* is decreased in **put()**; it is blocked if *empty* is zero;
- *full* is decreased in **get()**; it is blocked if *full* is zero.

## Nested Monitors - Bounded Buffer Program

```
const Max = 2
range Int = 0..Max
SEMAPHORE (N=0) = SEMA[N] ,
SEMA[v:Int]      = ( up -> SEMA[v+1]
                    | when (v>0) down ->SEMA[v-1] ) ,
SEMA[Max+1]      = ERROR.

BUFFER = (put -> empty.down -> full.up ->BUFFER
          | get -> full.down -> empty.up ->BUFFER) .

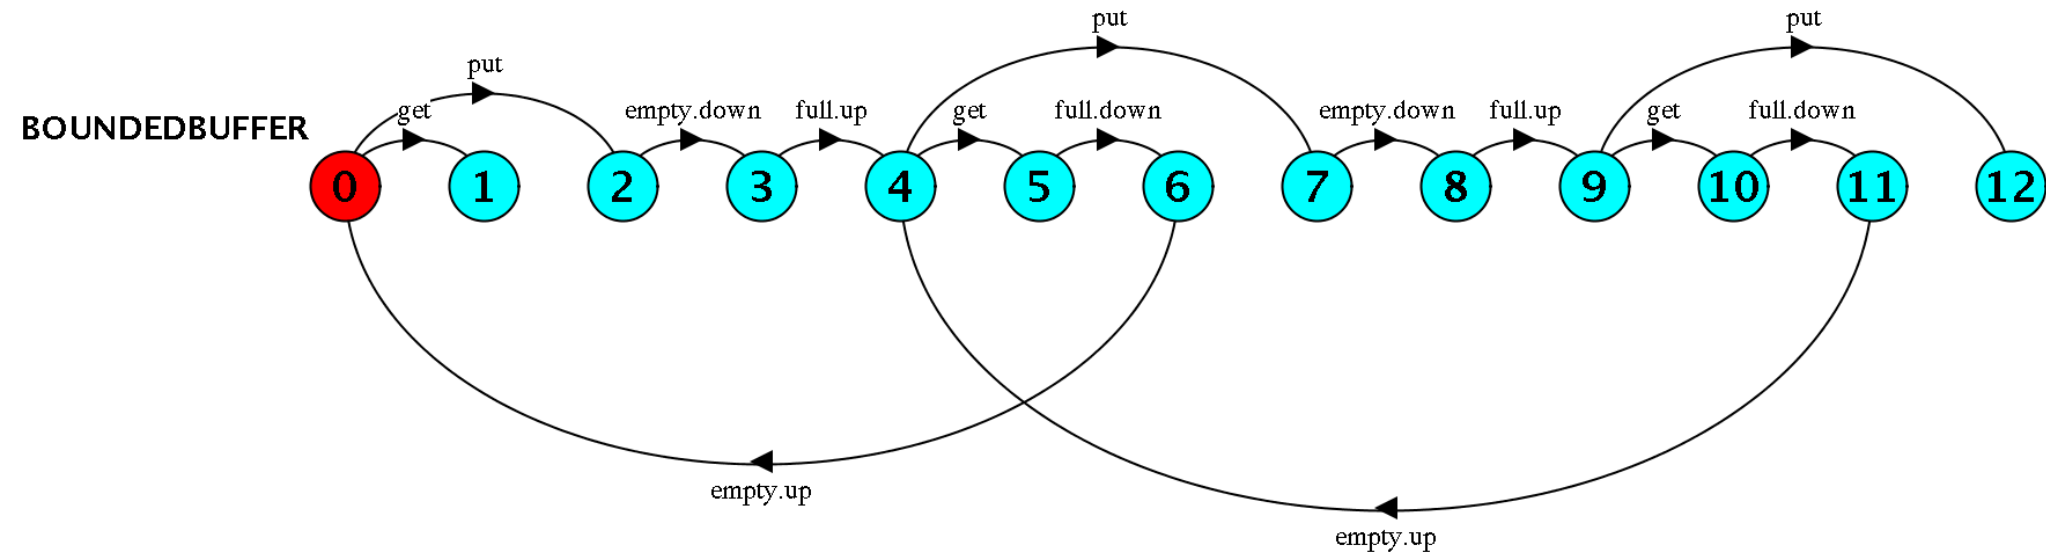
PRODUCER = (put -> PRODUCER) .
CONSUMER = (get -> CONSUMER) .

|| BOUNDEDBUFFER = (PRODUCER || BUFFER || CONSUMER
                    || empty:SEMAPHORE (2)
                    || full:SEMAPHORE (0)
                    ) .
```

*Does this behave  
as desired?*

# Nested Monitors - Bounded Buffer Program


---



## Nested Monitors - Bounded Buffer Program

---

```
synchronized public Object get()  
    throws InterruptedException{  
    full.down(); // if there is no item in Buffer  
    ...  
}
```



## Nested Monitors - Bounded Buffer Program

---

*LTSA* analysis predicts a possible **DEADLOCK**:

Composing

potential **DEADLOCK**

States Composed: 28 Transitions: 32 in 60ms

Trace to **DEADLOCK**:

**get**

The **Consumer** tries to **get** an item, but the buffer is **empty**. It **blocks** and **releases** the **lock** on the **semaphore full**. The **Producer** tries to **put** an item into the **buffer**, but also **blocks**.

This situation is known as the ***nested monitor problem***.

## Nested Monitors – Revised Bounded Buffer Program

---

The **deadlock** can be removed by ensuring that the **monitor lock** for the **buffer** is not **acquired** until *after* semaphores are decreased.

```
synchronized public E get() throws InterruptedException {
    full.down();
    E o =buf[out]; buf[out]=null; --count; out=(out+1)%size;
    empty.up();
    return (o);
}    // original

public E get() throws InterruptedException {
    full.down();
    synchronized(this) {
        E o =buf[out]; buf[out]=null; --count; out=(out+1)%size;
        empty.up();
        return (o);
    } // new
```

## Nested Monitors – Revised Bounded Buffer Program

---

The **deadlock** can be removed by ensuring that the **monitor lock** for the **buffer** is not **acquired** until *after* semaphores are decreased.

```
synchronized public void put(E o) throws InterruptedException{
    empty.down();
    buf[in] = o; ++count; in=(in+1)%size;
    full.up();
} // original

public void put(E o) throws InterruptedException {
    empty.down();
    synchronized(this) {
        buf[in] = o; ++count; in=(in+1)%size; }
    full.up();
} // new
```

## Nested Monitors – Revised Bounded Buffer Program

---


```
// original
BUFFER = (put -> empty.down ->full.up ->BUFFER
          |get -> full.down ->empty.up ->BUFFER) .
PRODUCER = (put -> PRODUCER) .
CONSUMER = (get -> CONSUMER) .

// New
BUFFER = (put -> BUFFER
          |get -> BUFFER ) .
PRODUCER = (empty.down->put->full.up->PRODUCER) .
CONSUMER = (full.down->get->empty.up->CONSUMER) .
```

The semaphore actions have been moved to the producer and consumer. This is exactly as in the implementation where the semaphore actions are *outside* the monitor—get and put.



# Nested Monitors – Revised Bounded Buffer Program



# Summary

---

## ◆ Concepts

- **monitors**: encapsulated data + access procedures  
mutual exclusion + condition synchronization
- nested monitors

## ◆ Model

- guarded actions

## ◆ Practice

- private data and synchronized methods in Java
- **wait()**, **notify()** and **notifyAll()** for condition synchronization
- single thread active in the monitor at a time